



# AlexNet and ZFNet

# Recap: convolutional networks

- Filter size, depth, stride, max-pooling, parameter sharing



# How this will go

- I'm going to explain what they did, and you guys will tell me why it made sense


# The dataset for AlexNet

- ImageNet has 15 million images in 22,000 categories
  - Old machine learning methods such as SVMs can't handle this well
- New GPU technology allowed for parallelizing Convolutional Neural Network (CNN) training
- They used 1.2 million samples and 1000 categories for training
- They downsampled all images to 256x256 pixels
  - Why do we need consistent sizes?


# The activation function

- They used ReLU's instead of sigmoid or tanh as their activation function


Sigmoid



TanH




ReLU



# Splitting across GPUs

- ~million training examples
- Used two GPUs, put half of neurons on each
  - GPUs communicated findings at third layer



# Local Response Normalization

- The predecessor to batch-norm; it is useful when using unbounded activation functions (such as relu which goes  $0-\infty$ )
- Batch-norm sets each hidden layer to 0 mean, unit variance
  - And adds two learnable parameters
- LRN favors the most activated neurons in their respective neighborhoods

$$b_{x,y}^i = a_{x,y}^i / (k + \alpha \sum_{j=\max(0, i-n/2)}^{j=\min(N-1, i+n/2)} a_{x,y}^j)^2)^\beta$$

where

$b_{x,y}^i$  – regularized output for kernel  $i$  at position  $x, y$

$a_{x,y}^i$  – source output of kernel  $i$  applied at position  $x, y$


$N$  – total number of kernels

$n$  – size of the normalization neighbourhood

$\alpha, \beta, k, (n)$  – hyperparameters

# The Architecture

- 5 convolutional layers, followed by 3 fully connected layers
- Why do convolutions get narrower/deeper further in network?
- What purpose do the three regular layers serve?
- What is the last layer's output?
- I thought the input was 256x256... Why's it 224x224?





# Data Augmentation to Reduce Overfitting

- Take several random 224x224 images from each 256x256 image
  - With their horizontal flips as well
- Then, modify the RGB values linearly using PCA and eigenvalues
  - This scales the color intensities up and down randomly
  - Why is this helpful?

# Dropout ( $p = 0.5$ ) to reduce Overfitting

- What is dropout?
- Why does it take twice as long to train?
- Why does it reduce overfitting?

# AlexNet Optimizer

- Momentum Gradient Descent
- Weight Decay

$$v_{i+1} := 0.9 \cdot v_i - 0.0005 \cdot \epsilon \cdot w_i - \epsilon \cdot \left\langle \frac{\partial L}{\partial w} \Big|_{w_i} \right\rangle_{D_i}$$

$$w_{i+1} := w_i + v_{i+1}$$

where  $i$  is the iteration index,  $v$  is the momentum variable,  $\epsilon$  is the learning rate, and  $\left\langle \frac{\partial L}{\partial w} \Big|_{w_i} \right\rangle_{D_i}$  is the average over the  $i$ th batch  $D_i$  of the derivative of the objective with respect to  $w$ , evaluated at  $w_i$ .

# AlexNet Weight Initialization

- Biases all initialized to 1, so the ReLU activation function will work its magic
  - Why does  $>0$  bias help ReLU?
- Weights initialized to mean 0 and variance 0.01
  - Why aren't the weights all the same values?

# Results

- This paper is the reason that deep learning is such a hot discipline right now
- They got 16.4% error compared to 2nd place's 26.1%
  - Using an unconventional strategy!
- Subsequent years saw results go as low as 5% using spinoffs of AlexNet
- Drawbacks:
  - Only recognize from 1000 categories
  - No explanation of context

## Cool! Let's move on to the second paper, ZF-Net

- ZF stands for Matthew Zeiler and Rob Fergus, the NYU researchers who made it
- Won the ImageNet challenge the year after AlexNet using a similar architecture
- Paper's official name is Visualizing and Understanding Deep Neural Networks
  - Made methods to see what a convolutional net 'knows'

# Post-AlexNet

- Improvement of neural networks little more than trial-and-error
  - But trials take a week or two to train!
- ZF-Net introduced visualization techniques that show what neurons are stimulated by specific inputs
  - “Deconvnet”

# Deconvnet

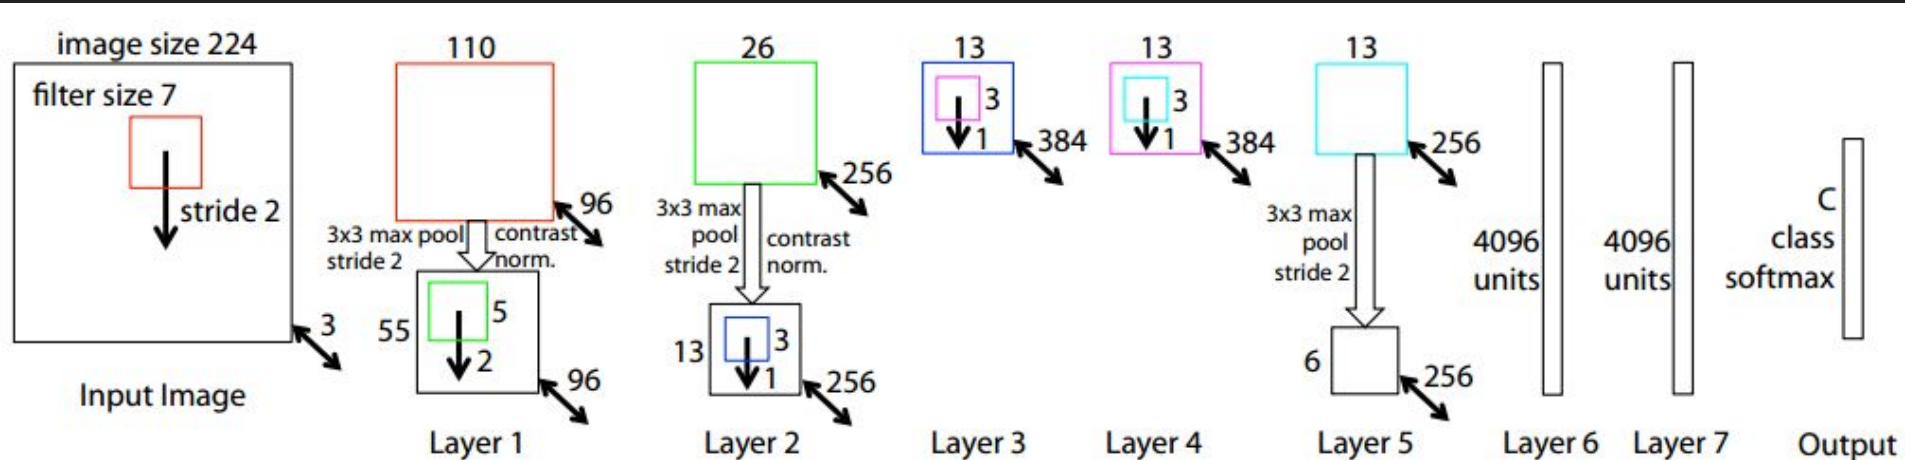
- At every layer, we map the activated neurons back to the pixels
- **Unpooling:** How do we reverse the max pooling layer?
  - Loss of information?
- **Rectification:** ReLU keeps everything positive
- **Filtering:** Which neurons from the previous layer contributed to neuron\_i in the current layer

The reconstruction we obtain of the pixels is not complete! The pixels are instead shown weighted by their contribution



# ZF-Net Architecture


- Their architecture started as the AlexNet architecture
- Looked at deconvnet results on AlexNet
- Changed architecture as necessary




# AlexNet architecture deconvnet results

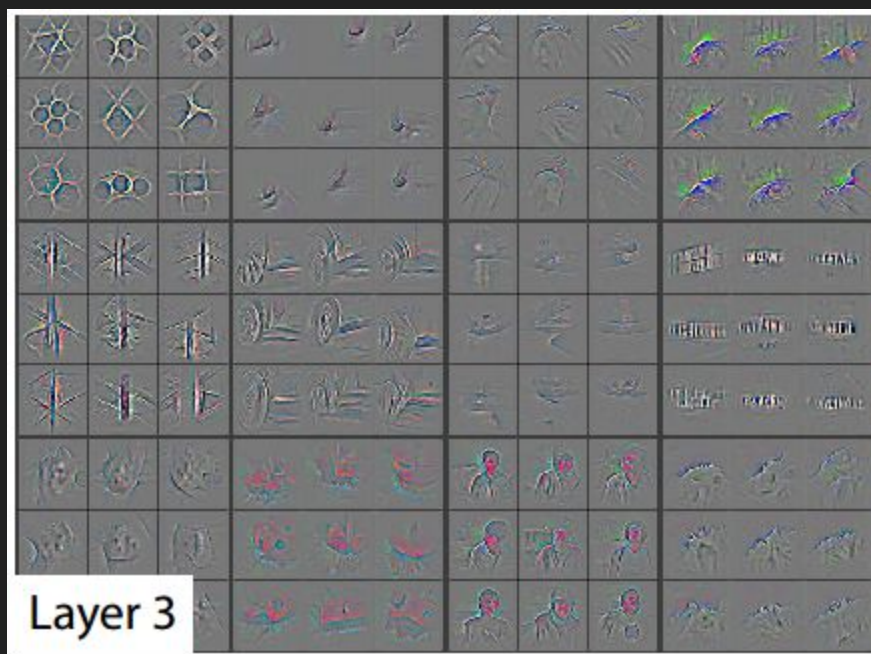


Layer 1




Layer 2

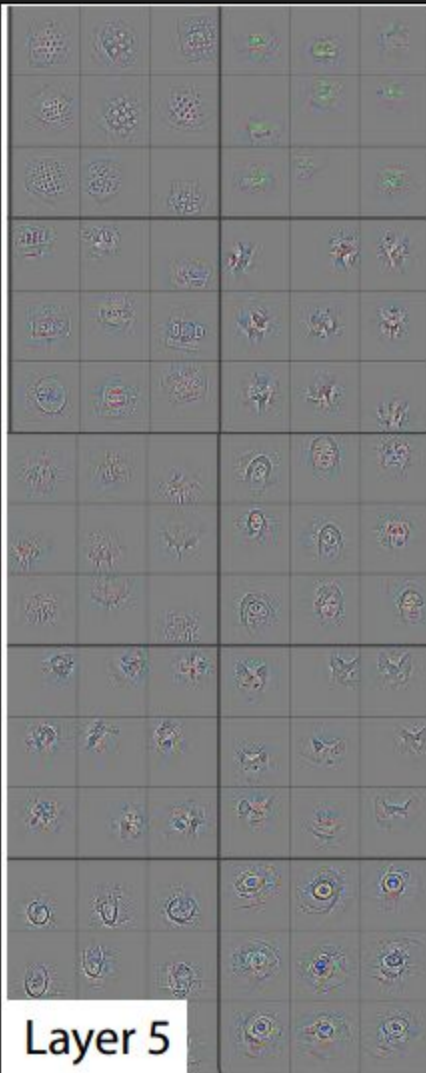










Layer 4



Layer 5



# Reminder of CNN method




# Changes based on deconvolutions

- First Layer responds only to stark differences in intensity
  - Little information passed on about medium intensity
  - Second Layer then has trouble putting together complicated patterns
- Old architecture did 11x11 squares and stride of 4 for first layer
- New architecture does 7x7 squares and stride of 2
  - Why does this fix the problem?




c, e are 1st, 2nd layers of ZF-Net; b, d are same for AlexNet

- ZF-Net's features more precise by second layer, due to first layer having more detail




(b)



(c)



(d)



(e)

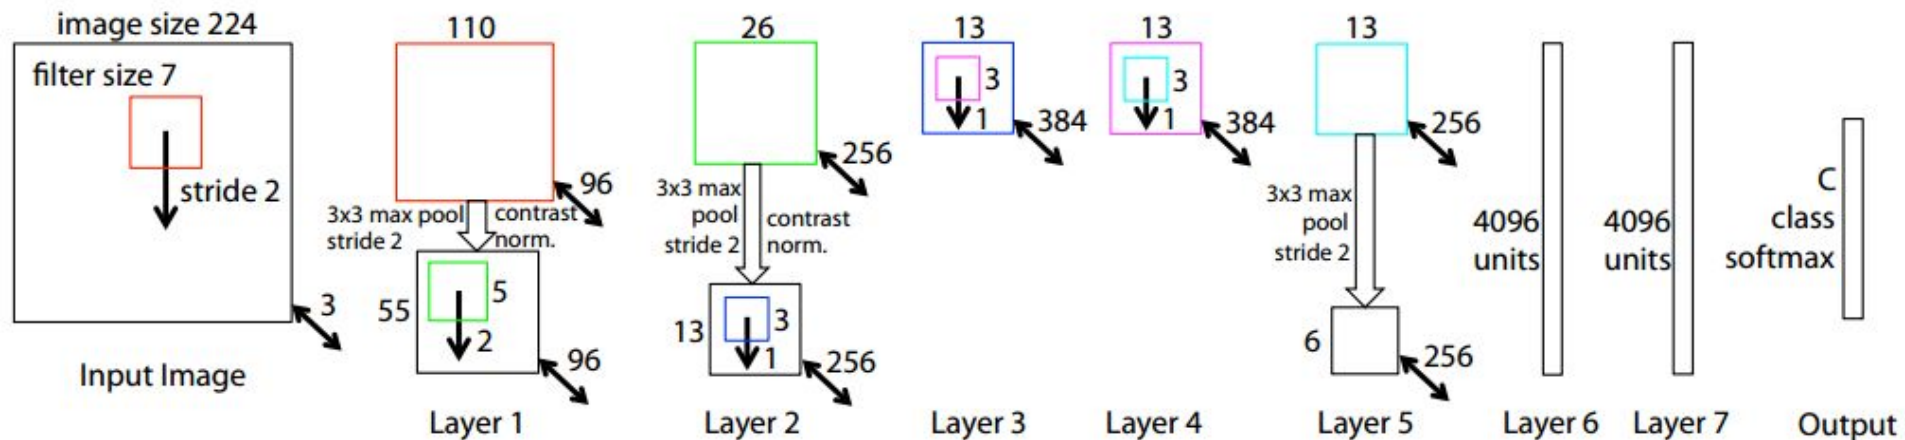
# Occlusion Sensitivity

- Is the model actually learning on the important section of the image or the environment?
- Randomly covered parts of the input with a gray square
- What if we cover the part that the convnet is sensitive to?



# Generalization

- Use a brand new dataset, but *keep the old neural network*
- Only allow learning on the last layer's parameters
- How does this make sense?
- Achieved 15% error using this method on brand new datasets!



# Next week

- [Dynamic Pooling and Unfolding Recursive Autoencoders for Paraphrase Detection](#)
- Their code is available [here](#) (but you don't need to look at it for the paper)